

A Compact Bytecode Format for **JavaScriptCore**

Tadeu Zagallo

Apple Inc.



webkit.org



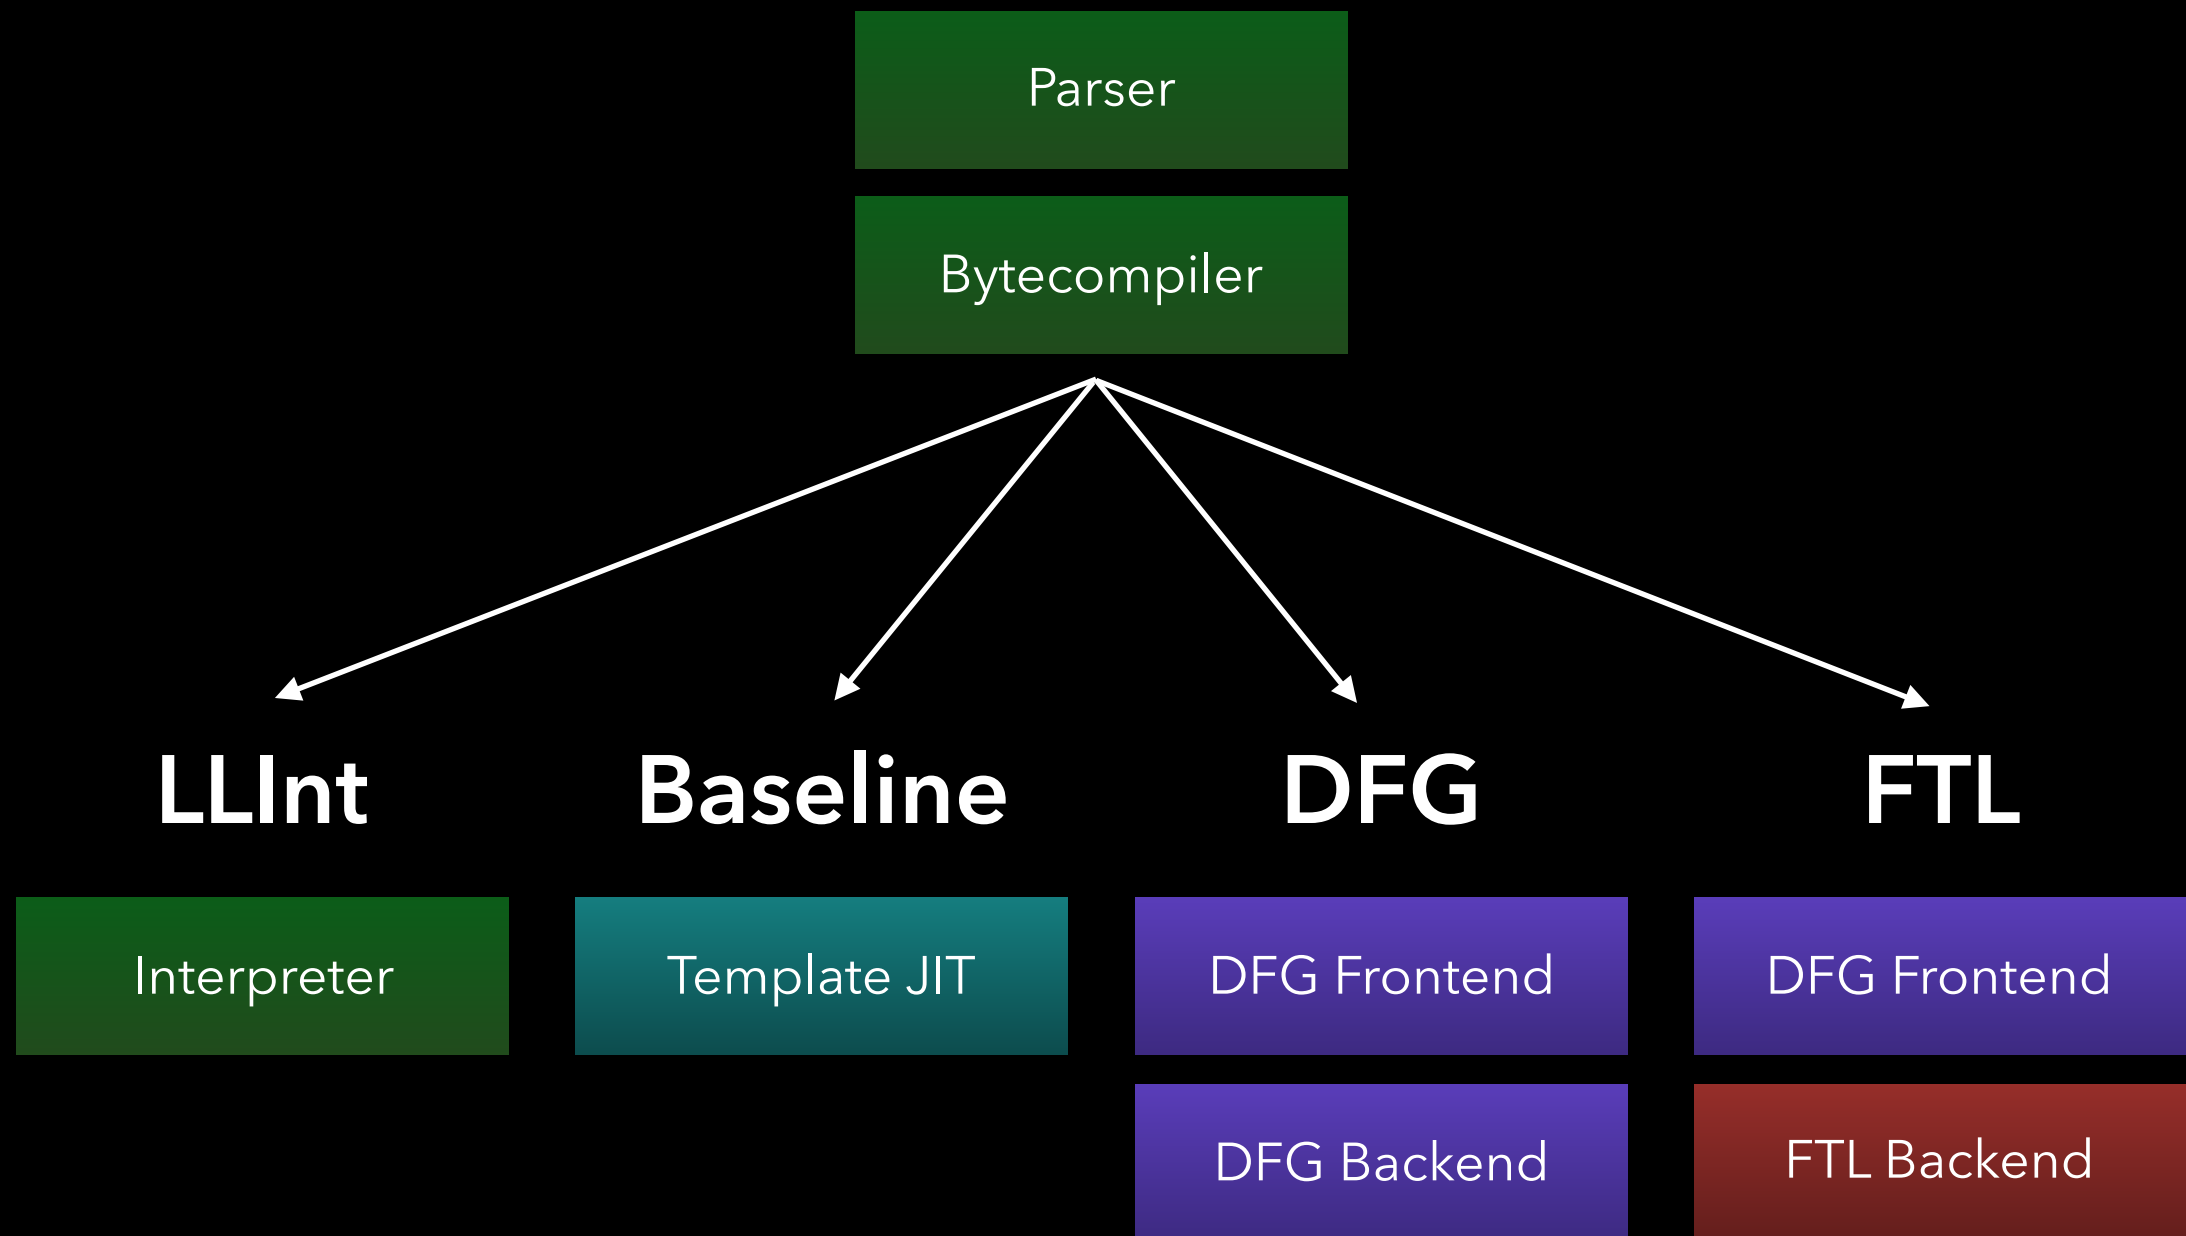
Safari

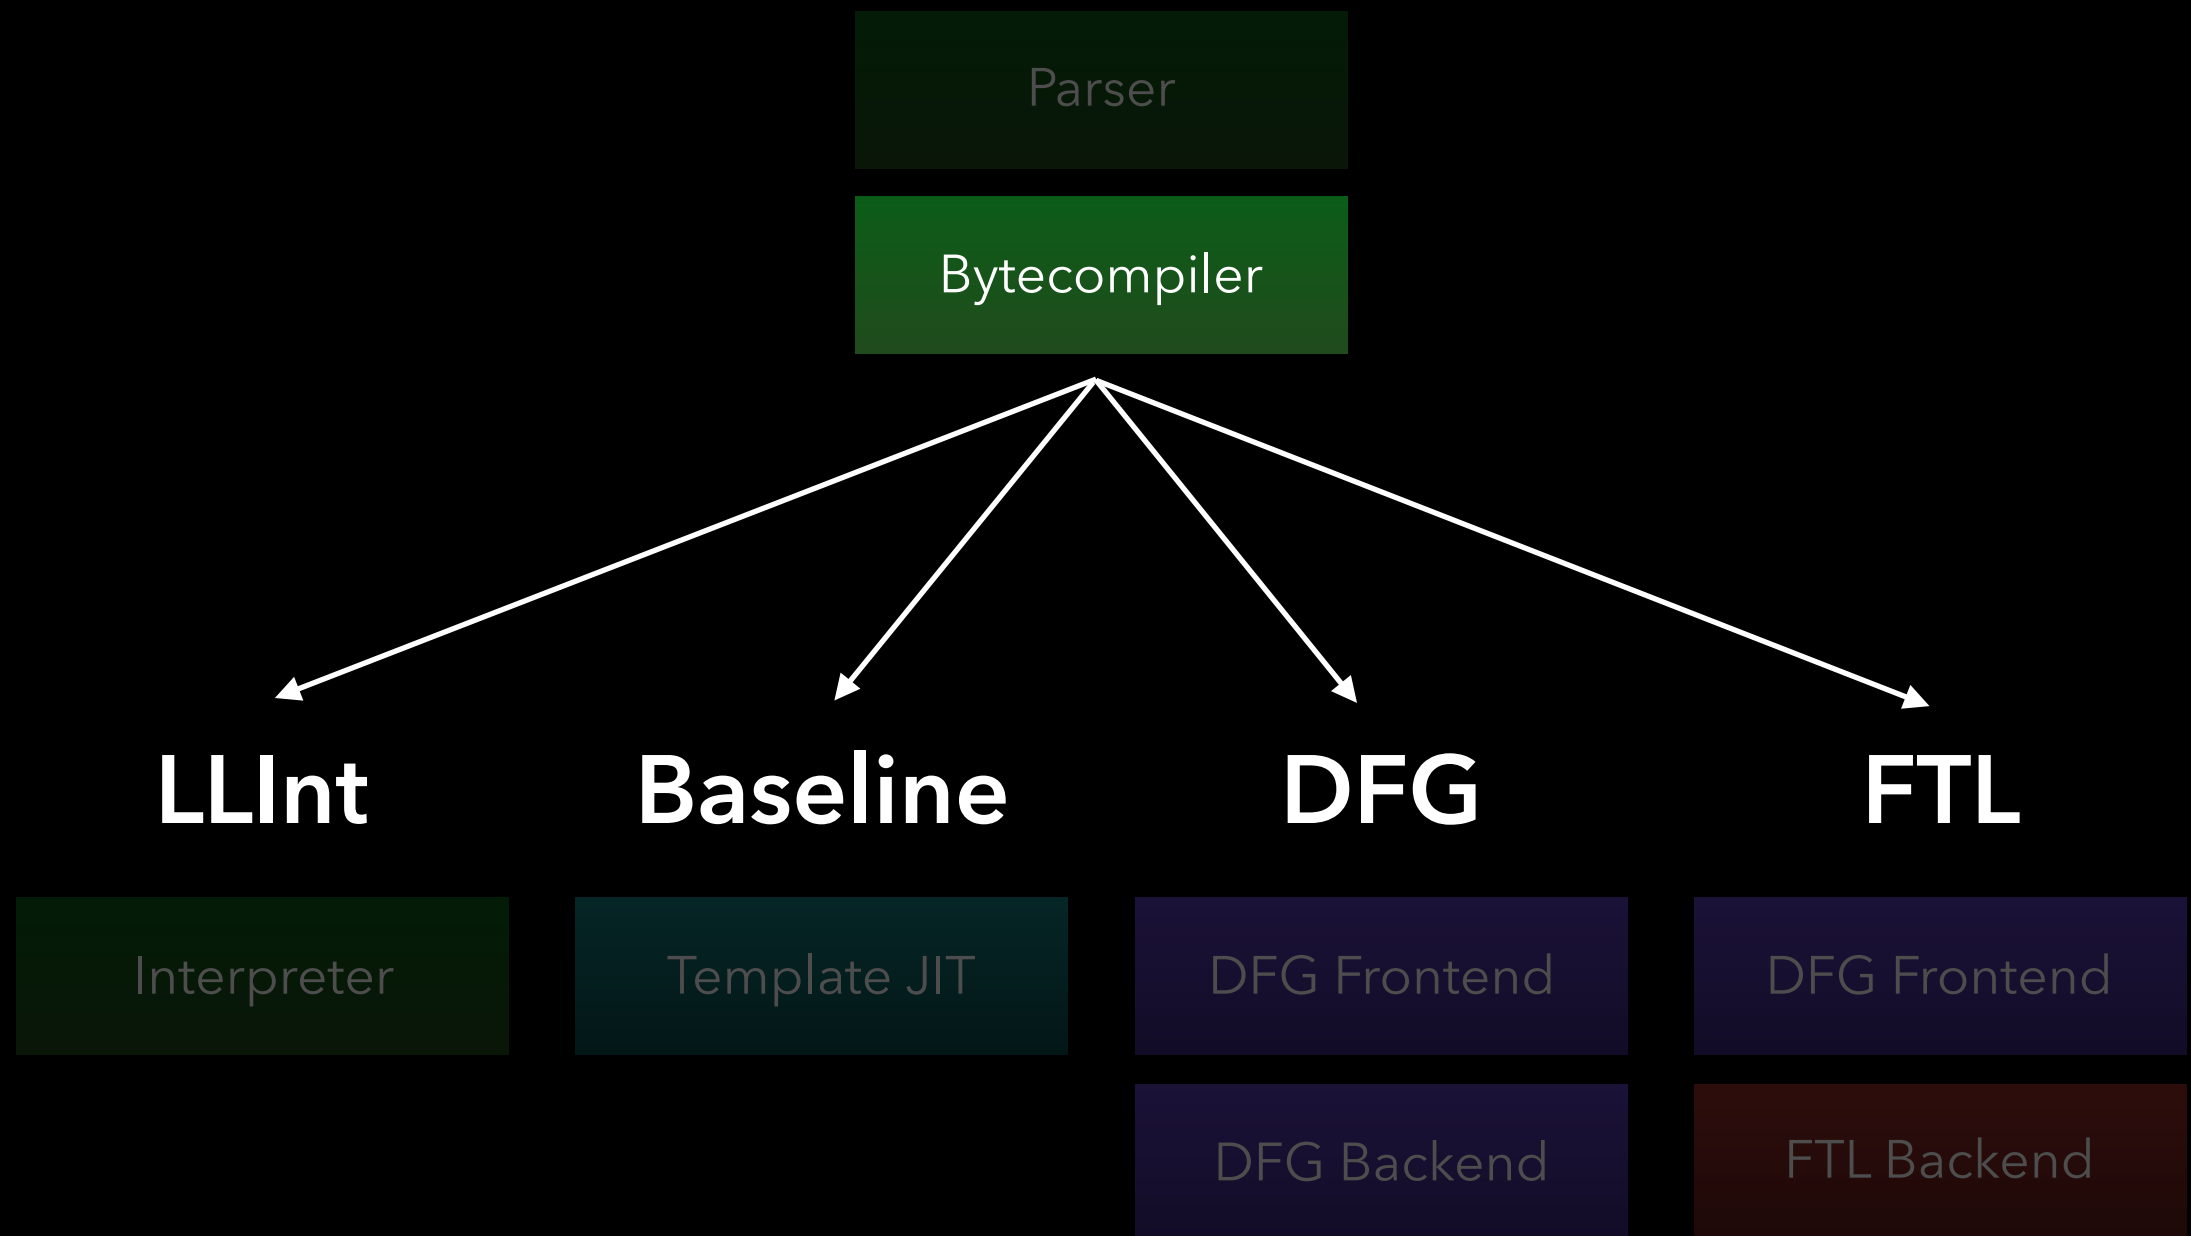
Agenda

- High level overview
- Old bytecode format
- New bytecode format
- Memory comparison
- Type safety improvements

Agenda

- High level overview
- Old bytecode format
- New bytecode format
- Memory comparison
- Type safety improvements





Bytecode Goals

- Memory efficiency
- Cacheable

Bytecode

```
// double.js  
function double(a) {  
    return a + a;  
}  
double(2);
```

```
$ jsc -d double.js
```

Bytecode

```
[ 0] enter
[ 1] get_scope      loc4
[ 3] mov            loc5, loc4
[ 6] check_traps
[ 7] add            loc7, arg1, arg1,
                  OperandTypes(126, 126)
[13] ret            loc7
```

Agenda

- High level overview
- Old bytecode format
- New bytecode format
- Memory comparison
- Type safety improvements

Old Bytecode Format

- Used too much memory
- The instruction stream was writable
- It had optimizations that were no longer beneficial

Old Bytecode Format

- Unlinked Instructions
 - Compact
 - Optimized for storage
- Linked Instructions
 - Inflated
 - Optimized for execution

Unlinked Instruction

1 byte

1 byte

1 byte

1 byte

2 bytes

op_add

0x1A

dst

0xF8

lhs

0x01

rhs

0x01

operandTypes

0xFEFE

Linked Instruction

8 bytes

8 bytes

8 bytes

8 bytes

8 bytes

<i>op_add</i>	<i>dst</i>	<i>lhs</i>	<i>rhs</i>	<i>arithProfile</i>
0x0000000010003240	0xFFFFFFFFFFFFFFFF8	0x0000000000000001	0x0000000000000001	0x00000000100039D8

Execution

- Direct threading
- Inline caching

Execution

- offlineasm overview
- Direct threading
- Inline caching

Execution

- offlineasm overview
- Direct threading
- Inline caching

offlineasm

```
macro load(tmp, getter)
    getter(tmp)
    loadi [tmp], tmp
end
```

```
_label:
    load(t0, macro(tmp) move 42, tmp end)
```

offlineasm

```
macro load(tmp, getter)
    getter(tmp)
    loadi [tmp], tmp
end

_label:
    load(t0, macro(tmp) move 42, tmp end)
```

Temporary registers: t0-t5

offlineasm

```
macro load(tmp, getter)
    getter(tmp)
    loadi [tmp], tmp
end
```

```
_label:
    load(t0, macro(tmp) move 42, tmp end)
```

Instruction suffixes

- b for byte
- i for 32-bit
- p for pointer
- h for 16-bit
- q for 64-bit

offlineasm

```
macro load(tmp, getter)
  getter(tmp)
  loadi [tmp], tmp
end
```

```
_label:
  load(t0, macro(tmp) move 42, tmp end)
```

Macros are lambda expressions that take zero or more arguments and return code

offlineasm

```
macro load(tmp, getter)
    getter(tmp)
    loadi [tmp], tmp
end

_label:
    load(t0, macro(tmp) move 42, tmp end)
```

Macros may be anonymous

offlineasm

```
macro load(tmp, getter)
    getter(tmp)
    loadi [tmp], tmp
end
```

```
_label:
    load(t0, macro(tmp) move 42, tmp end)
```

And macros can also be passed as arguments to other macros

Execution

- offlineasm overview
- Direct threading
- Inline caching

Direct Threading

8 bytes		8 bytes		8 bytes		8 bytes	
...	<i>op_mov</i> 0x000010011080	<i>dst</i> 0xFFFFFFFFFFFA	<i>src</i> 0xFFFFFFFFFFFB	<i>op_add</i> 0x000010003240	...		

↑
PC

```
macro dispatch(instructionSize)
    addp instructionSize * PtrSize, PC
    jmp [PC]
end
```

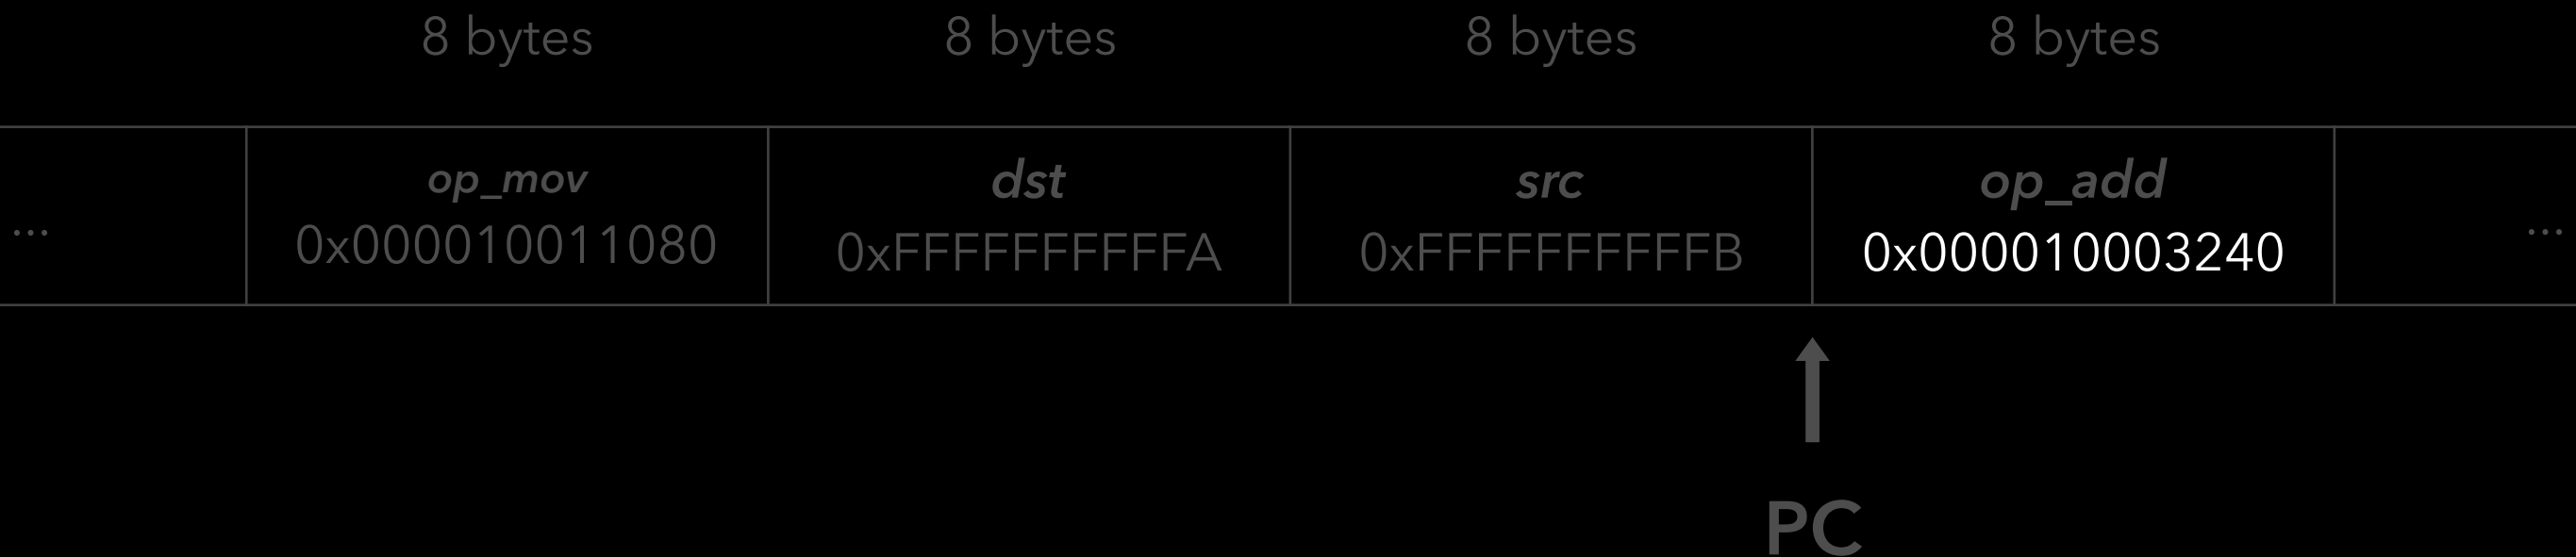
Direct Threading

8 bytes		8 bytes		8 bytes		8 bytes	
...	<i>op_mov</i> 0x000010011080	<i>dst</i> 0xFFFFFFFFFFFA	<i>src</i> 0xFFFFFFFFFFFB	<i>op_add</i> 0x000010003240	...		

↑
PC

```
macro dispatch(instructionSize)
    addp instructionSize * PtrSize, PC
    jmp [PC]
end
```

Direct Threading



```
macro dispatch(instructionSize)
    addp instructionSize * PtrSize, PC
    jmp [PC]
end
```

Execution

- offlineasm overview
- Direct threading
- Inline caching

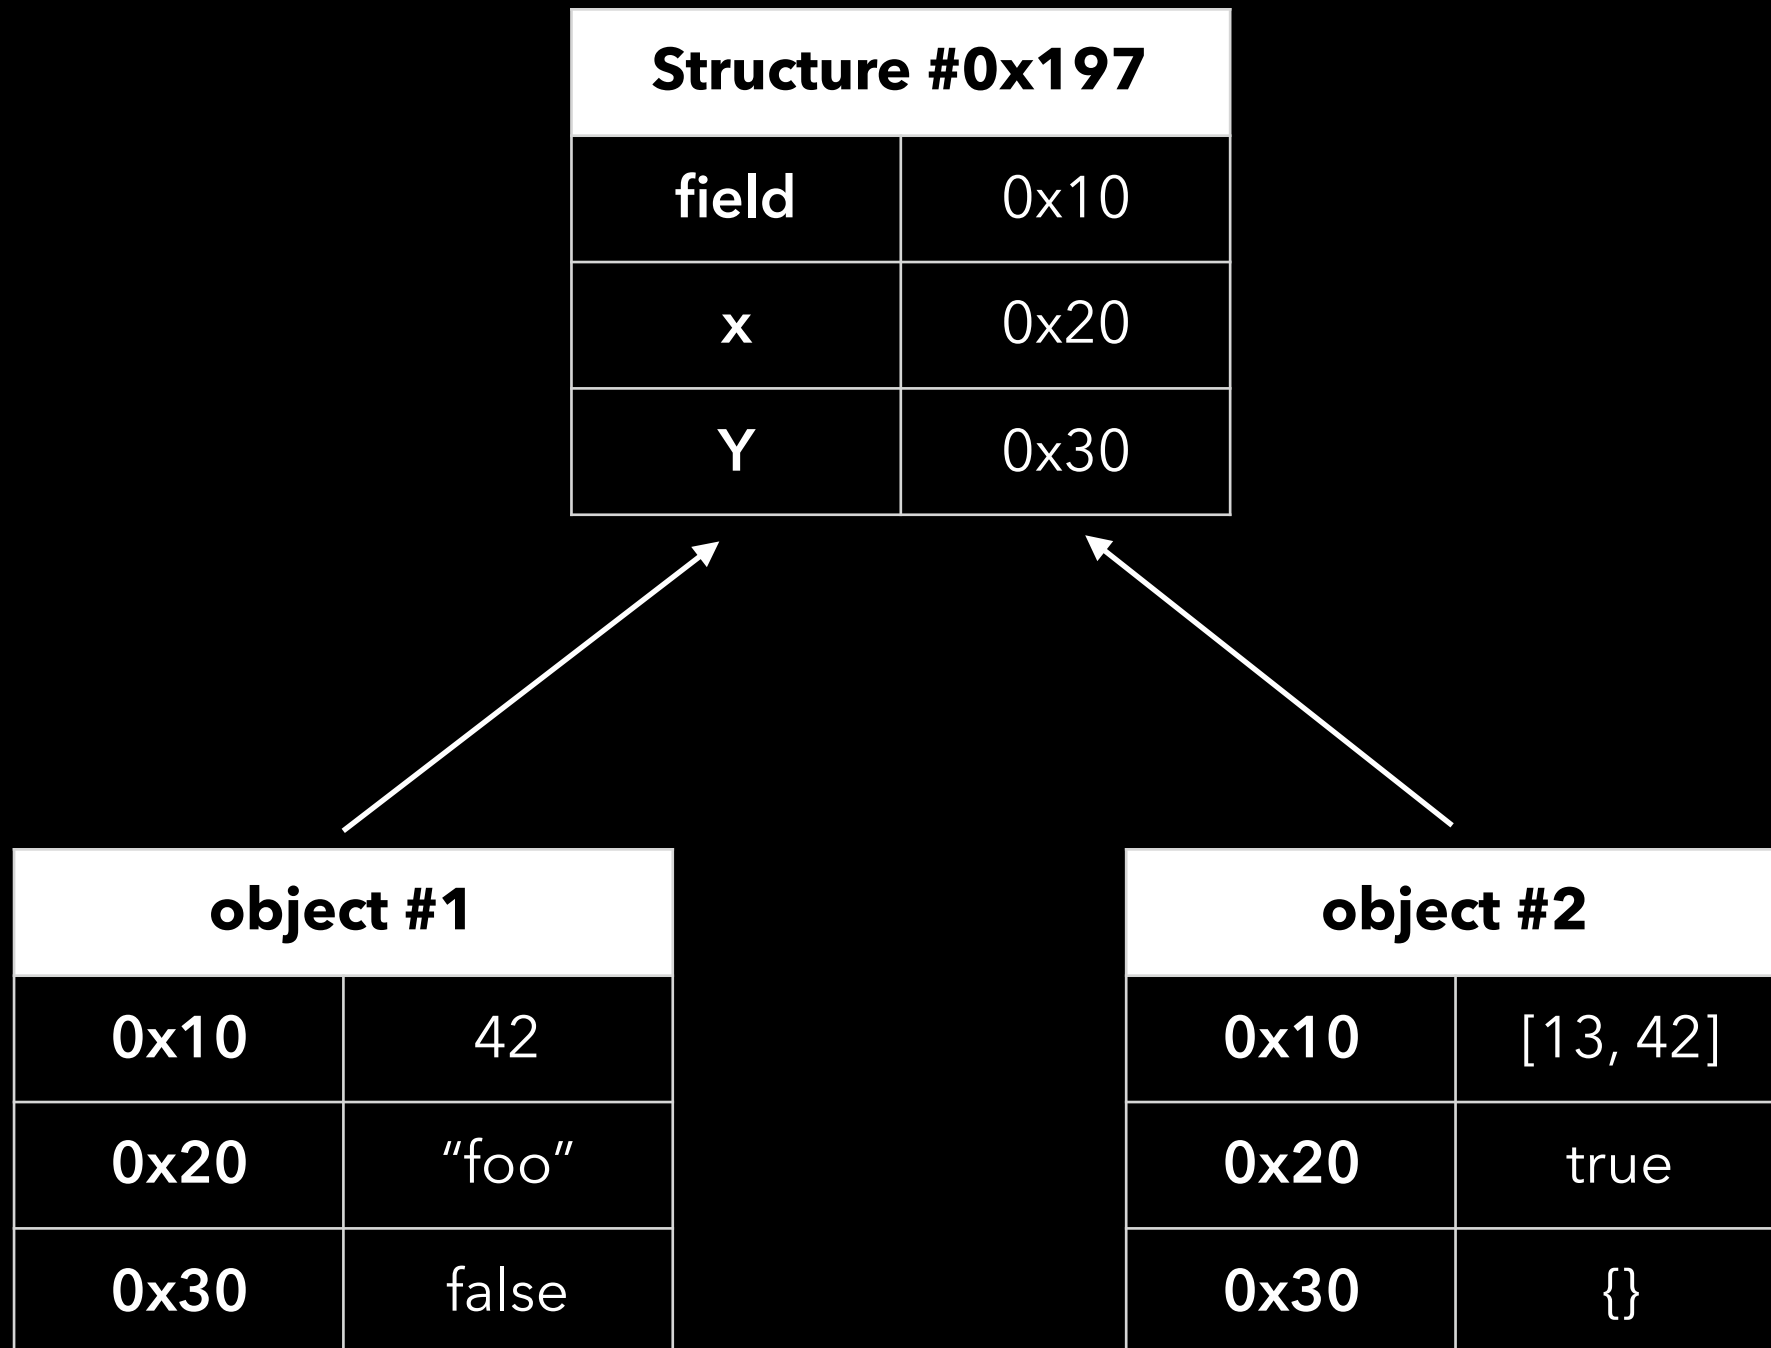
Inline Caching

`object.field`



`get_by_id object, field`

Inline Caching



Inline Caching

object.field



get_by_id object, field, 0, 0

Structure ID

Offset

object #1	
0x10	42
0x20	"foo"
0x30	false



Structure #0x197	
field	0x10
x	0x20
Y	0x30

object.field



get_by_id object, field, 0, 0

object #1	
0x10	42
0x20	"foo"
0x30	false



Structure #0x197	
field	0x10
x	0x20
y	0x30

object.field



get_by_id object, field, 0x197, 0x10

Agenda

- High level overview
- Old bytecode format
- New bytecode format
- Memory comparison
- Type safety improvements

New Bytecode

- Compact
 - No separate linked format
 - Multiple encoding sizes
- Cacheable
 - No runtime values
 - Read-only instruction stream

Narrow Instructions

1 byte

1 byte

1 byte

1 byte

1 byte

1 byte

op_add

0x1A

dst

0xF8

lhs

0x01

rhs

0x01

operandTypes

0xFE

metadataID

0x00

Wide Instructions

(32-bit words)

1 byte

4 bytes

4 bytes

4 bytes

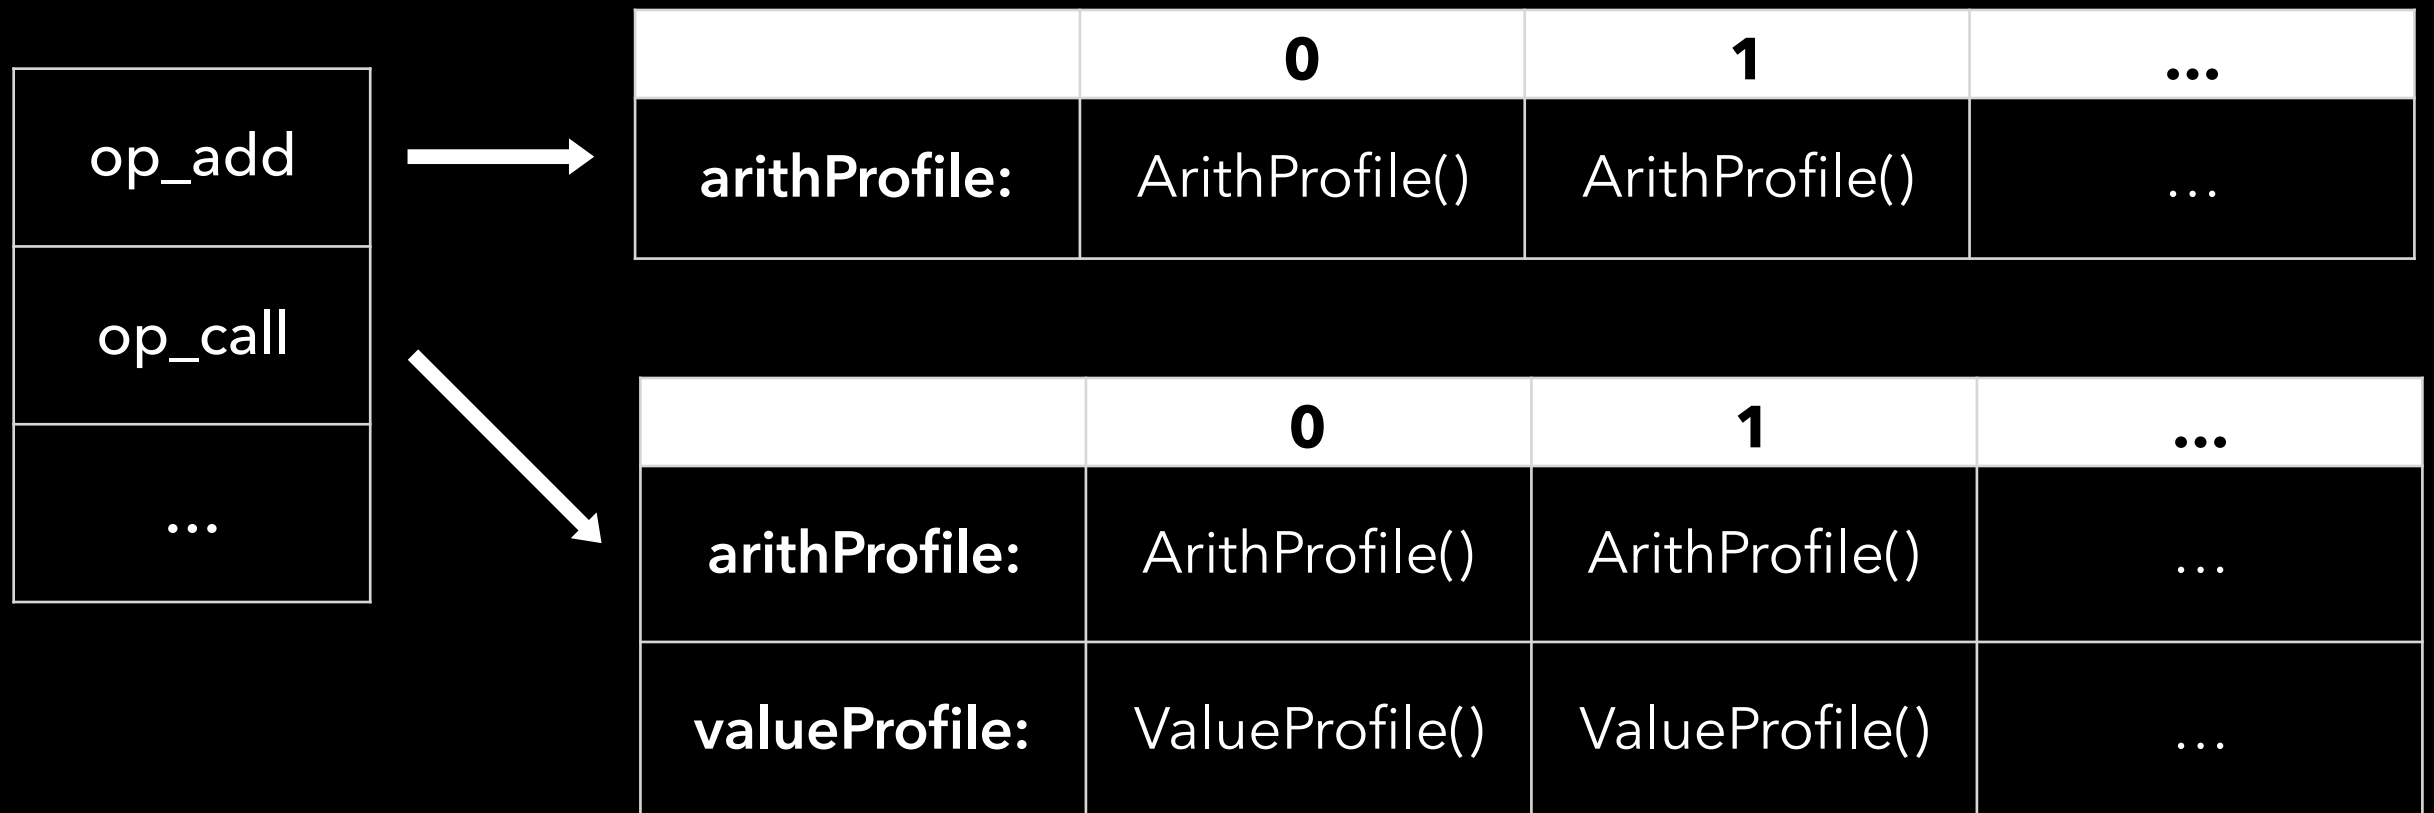
4 bytes

4 bytes

4 bytes

<i>op_wide</i>	<i>op_add</i>	<i>dst</i>	<i>lhs</i>	<i>rhs</i>	<i>operandTypes</i>	<i>metadataID</i>
0x01	0x0000001A	0xFFFFFFFF8	0x00000001	0x00000001	0xFFFFFFFFE	0x00010000

Metadata Table



Metadata Table

~200 opcodes × 8 bytes × ~23k tables

=

~36MB

Metadata Table

Header			Payload			
<i>0x0</i>	<i>0x4</i>	<i>...</i>	<i>0x100</i>	<i>0x110</i>	<i>0x120</i>	<i>...</i>
<i>op_add</i> <i>0x100</i>	<i>op_call</i> <i>0x120</i>	<i>...</i>	OpAdd::Metadata[0]	OpAdd::Metadata[1]	OpCall::Metadata[0]	<i>...</i>

- Allocate the whole table as a single chunk of memory
- Only allocate space for opcodes that have metadata
- Change the header from pointer to unsigned offset

Execution

- Indirect threading
- Inline caching
- Wide instruction execution

Execution

- Indirect threading
- Inline caching
- Wide instruction execution

Indirect Threading

```
macro dispatch(instructionSize)
    addp instructionSize * PtrSize, PC
    jmp [PC]
end
```

Indirect Threading

```
macro dispatch(instructionSize)
    addp instructionSize, PC
    loadb [PC], t0
    leap _g_opcodeMap, t1
    jmp [t1, t0, PtrSize]
end
```

Execution

- Indirect threading
- Inline caching
- Wide instruction execution

Inline Caching



Execution

- Indirect threading
- Inline caching
- Wide instruction execution

Wide Instruction Execution

```
macro dispatch(instructionSize)
    addp instructionSize, PC
    loadb [PC], t0
    leap _g_opcodeMap, t1
    jmp [t1, t0, PtrSize]
end
```

Wide Instruction Execution

```
macro dispatch(instructionSize)
    addp instructionSize, PC
    loadb [PC], t0
    leap _g_opcodeMap, t1
    jmp [t1, t0, PtrSize]
end
```

```
_llint_op_wide:
    loadi 1[PC], t0
    leap _g_opcodeMapWide, t1
    jmp [t1, t0, PtrSize]
```

Wide Instruction Execution

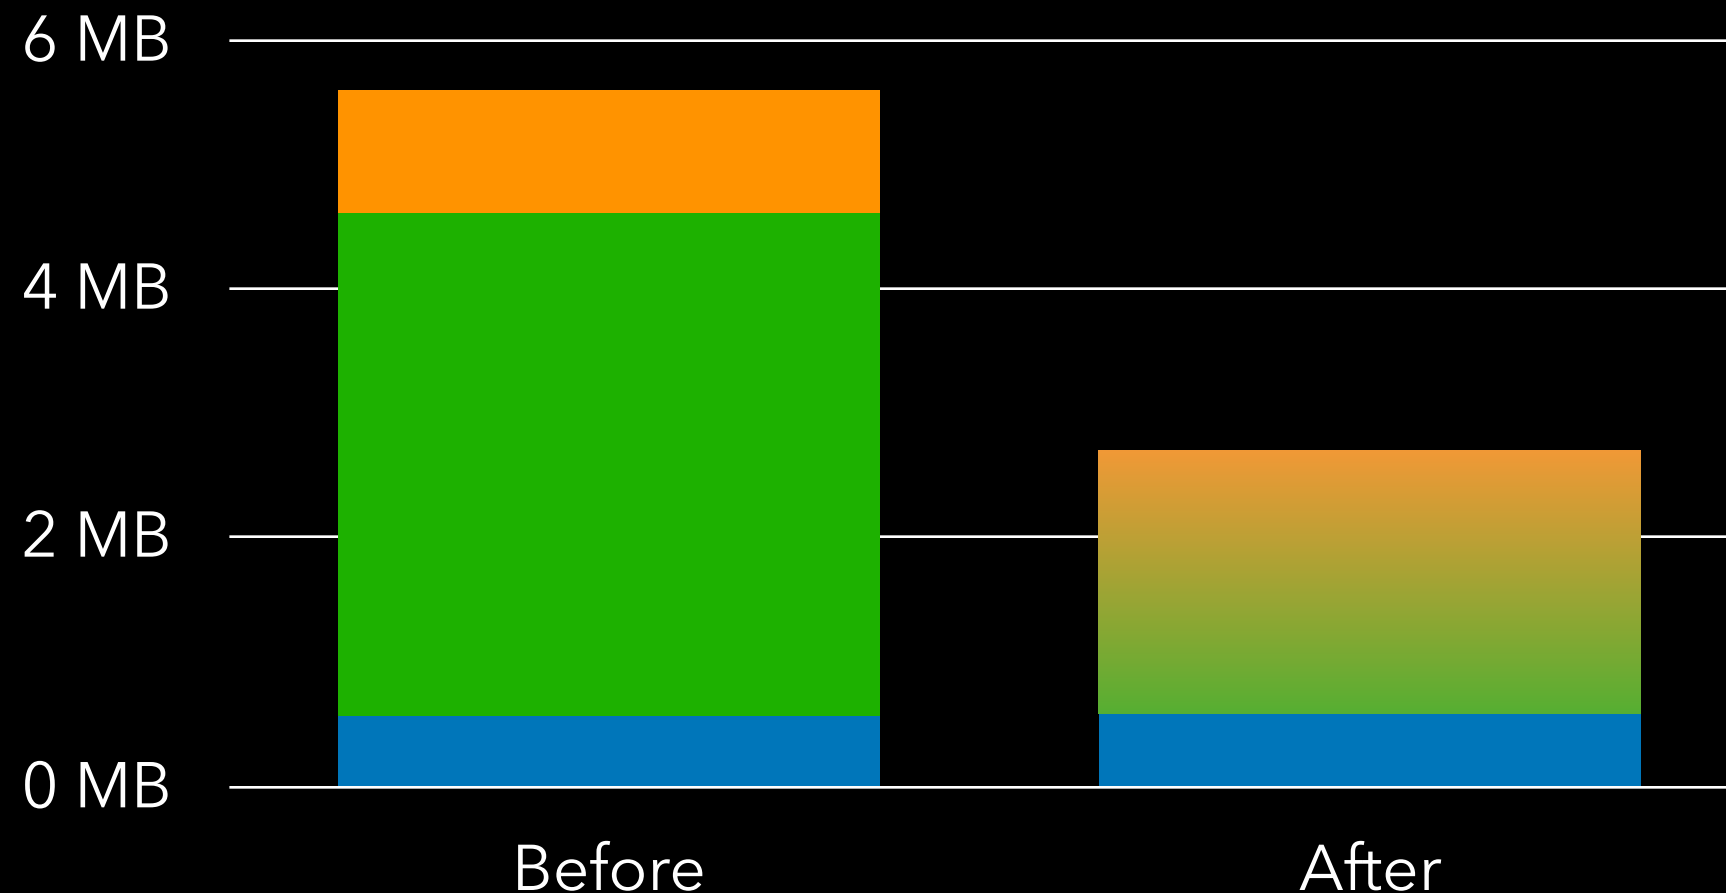
```
macro dispatch(instructionSize)
    addp instructionSize, PC
    loadb [PC], t0
    leap _g_opcodeMap, t1
    jmp [t1, t0, PtrSize]
end
```

```
_llint_op_wide:
    loadl 1[PC], t0
    leap _g_opcodeMapWide, t1
    jmp [t1, t0, PtrSize]
```

Agenda

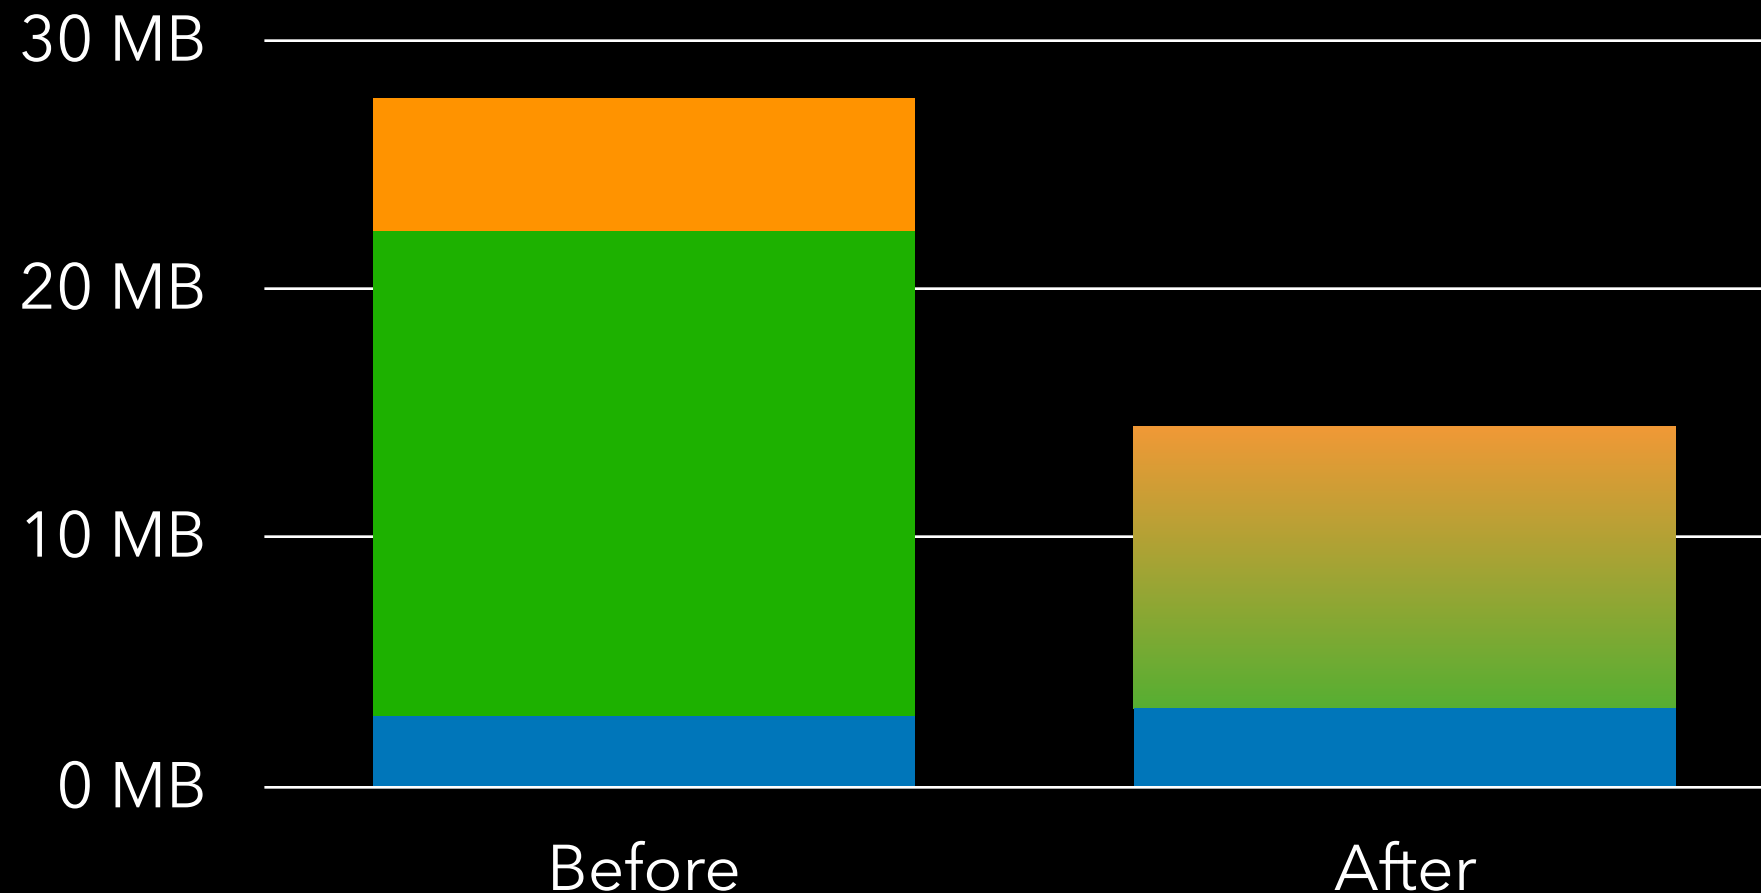
- High level overview
- Old bytecode format
- New bytecode format
- Memory comparison
- Type safety improvements

apple.com



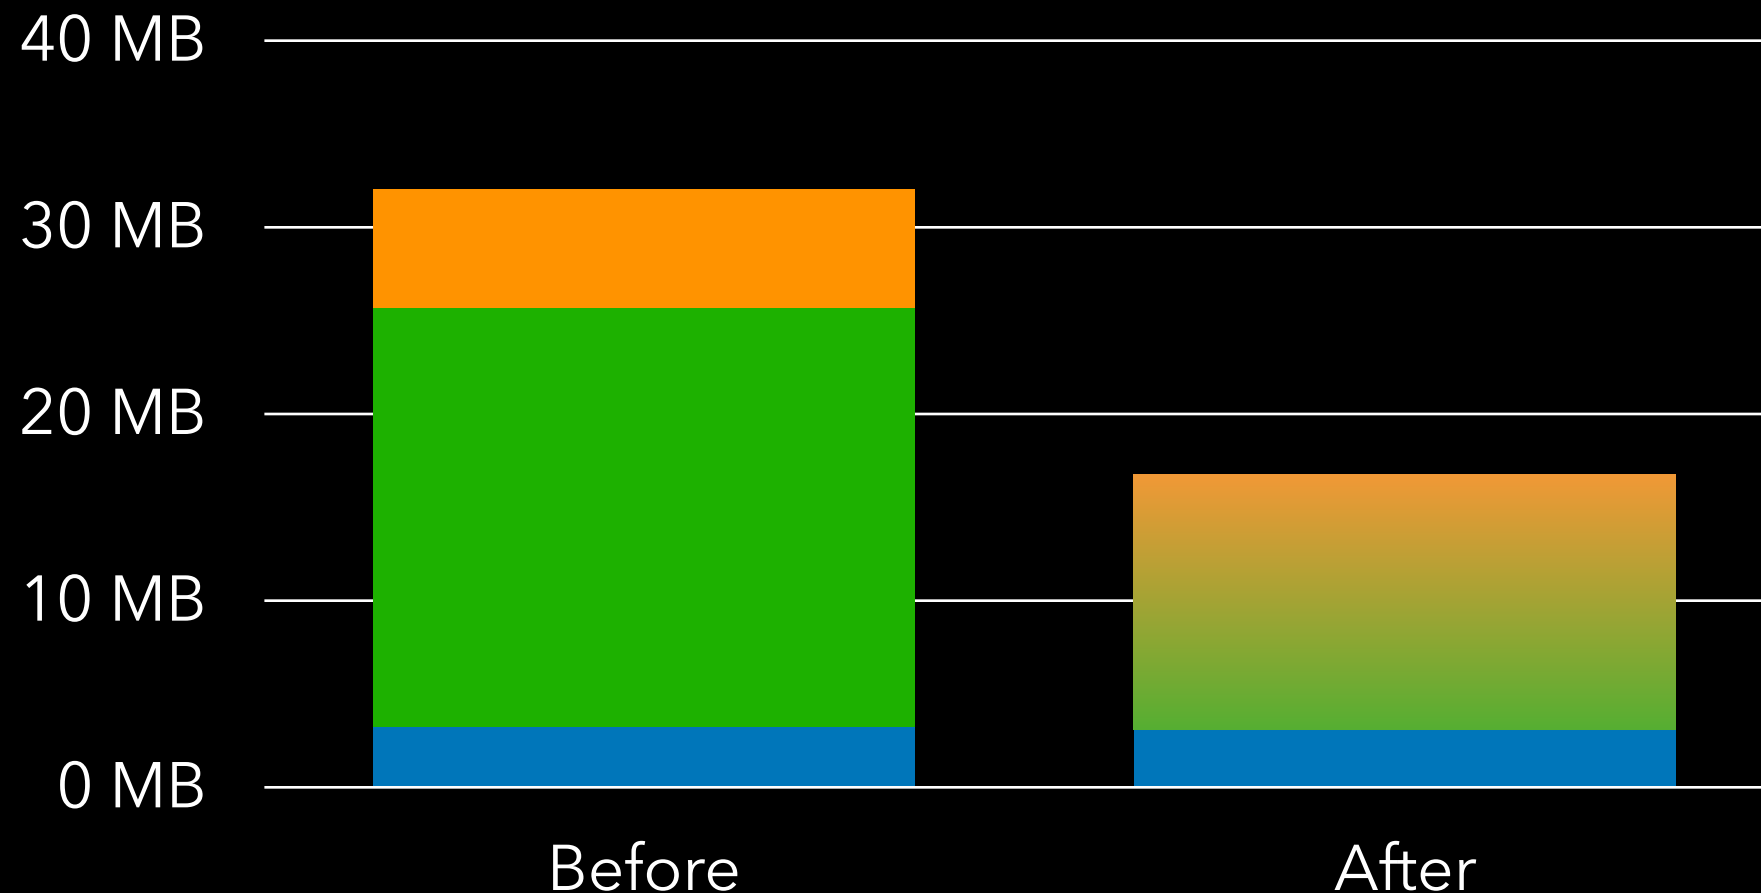
Description		Before	After	%
	Unlinked	0.55 MB	0.57 MB	+4%
	Linked	4.05 MB	2.14 MB	-57%
	Metadata	0.99 MB		
Total		5.60 MB	2.71 MB	-52%

reddit.com



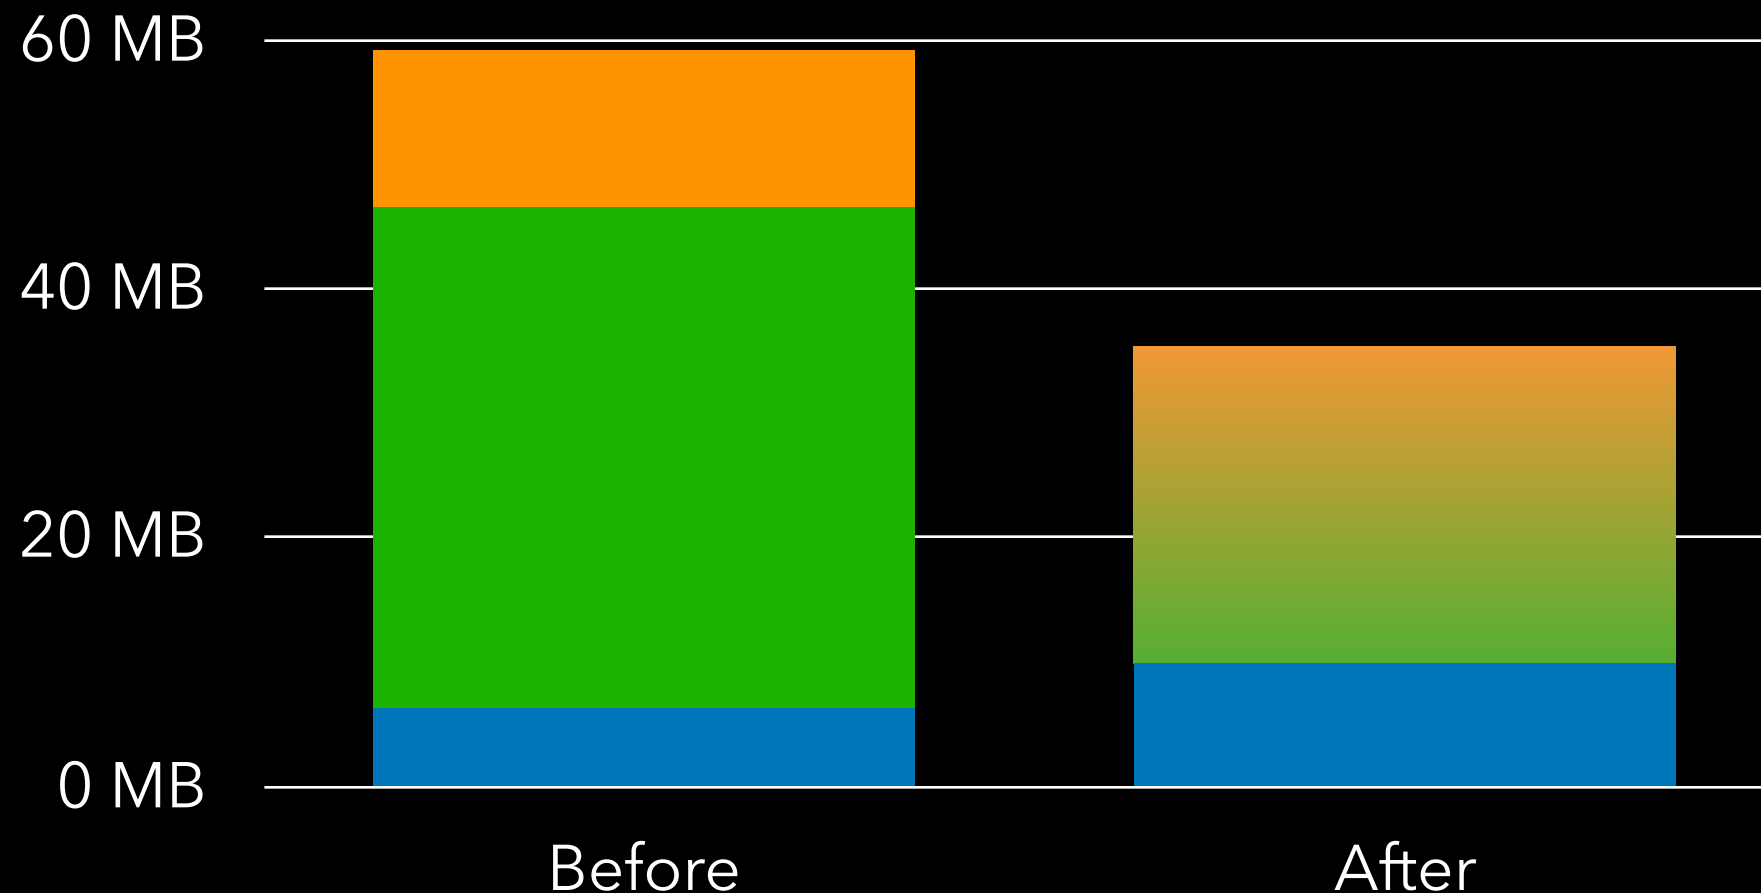
Description		Before	After	%
<div></div>	Unlinked	2.76 MB	3.08 MB	+12%
<div></div>	Linked	19.51 MB	11.37 MB	-54%
<div></div>	Metadata	5.34 MB		
Total		27.61 MB	14.45 MB	-48%

facebook.com



Description		Before	After	%
<div></div>	Unlinked	3.11 MB	2.99 MB	-4%
<div></div>	Linked	22.43 MB	13.66 MB	-52%
<div></div>	Metadata	6.51 MB		
Total		32.04 MB	16.65 MB	-48%

gmail.com



Description		Before	After	%
<div></div>	Unlinked	6.17 MB	9.89 MB	+60%
<div></div>	Linked	40.28 MB	25.51 MB	-52%
<div></div>	Metadata	12.75 MB		
Total		59.21 MB	35.40 MB	-40%

gmail.com

- More than 12k code blocks
- More than 830k instructions
- 270k wide instructions (33%)

Wide Instructions

(16-bit words)

1 byte

2 bytes

2 bytes

2 bytes

2 bytes

2 bytes

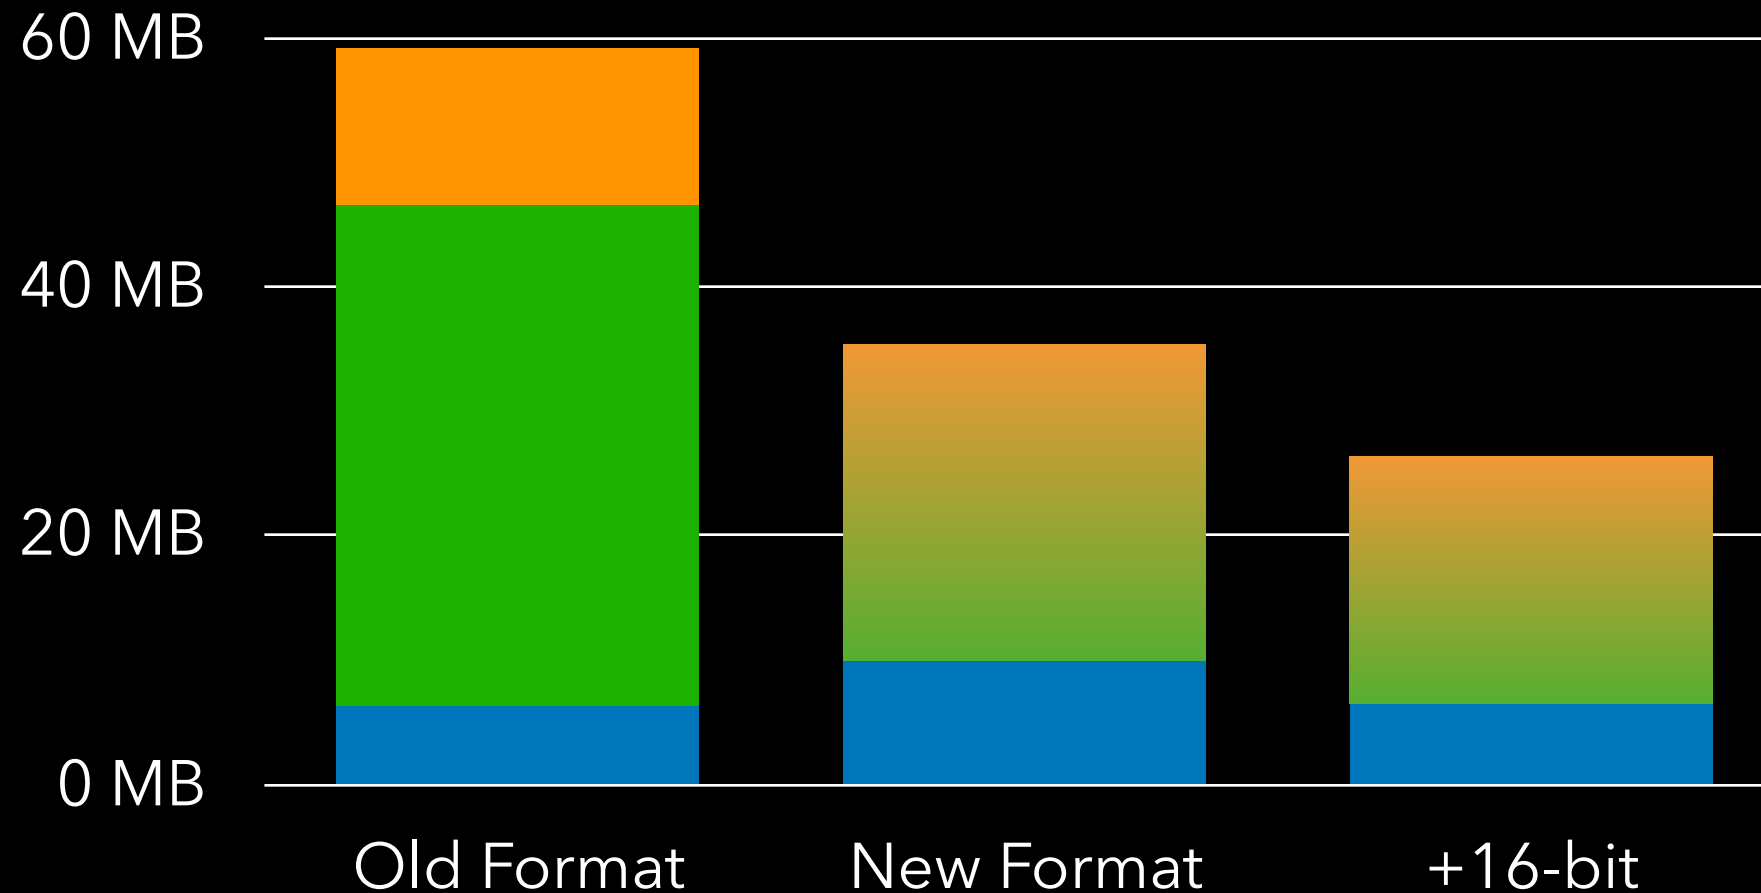
2 bytes




<i>op_wide16</i> 0x00	<i>op_add</i> 0x001A	<i>dst</i> 0xFFFF8	<i>lhs</i> 0x0001	<i>rhs</i> 0x0001	<i>operandTypes</i> 0xFEFE	<i>metadataID</i> 0x0100
--------------------------	-------------------------	-----------------------	----------------------	----------------------	-------------------------------	-----------------------------

Metadata Table

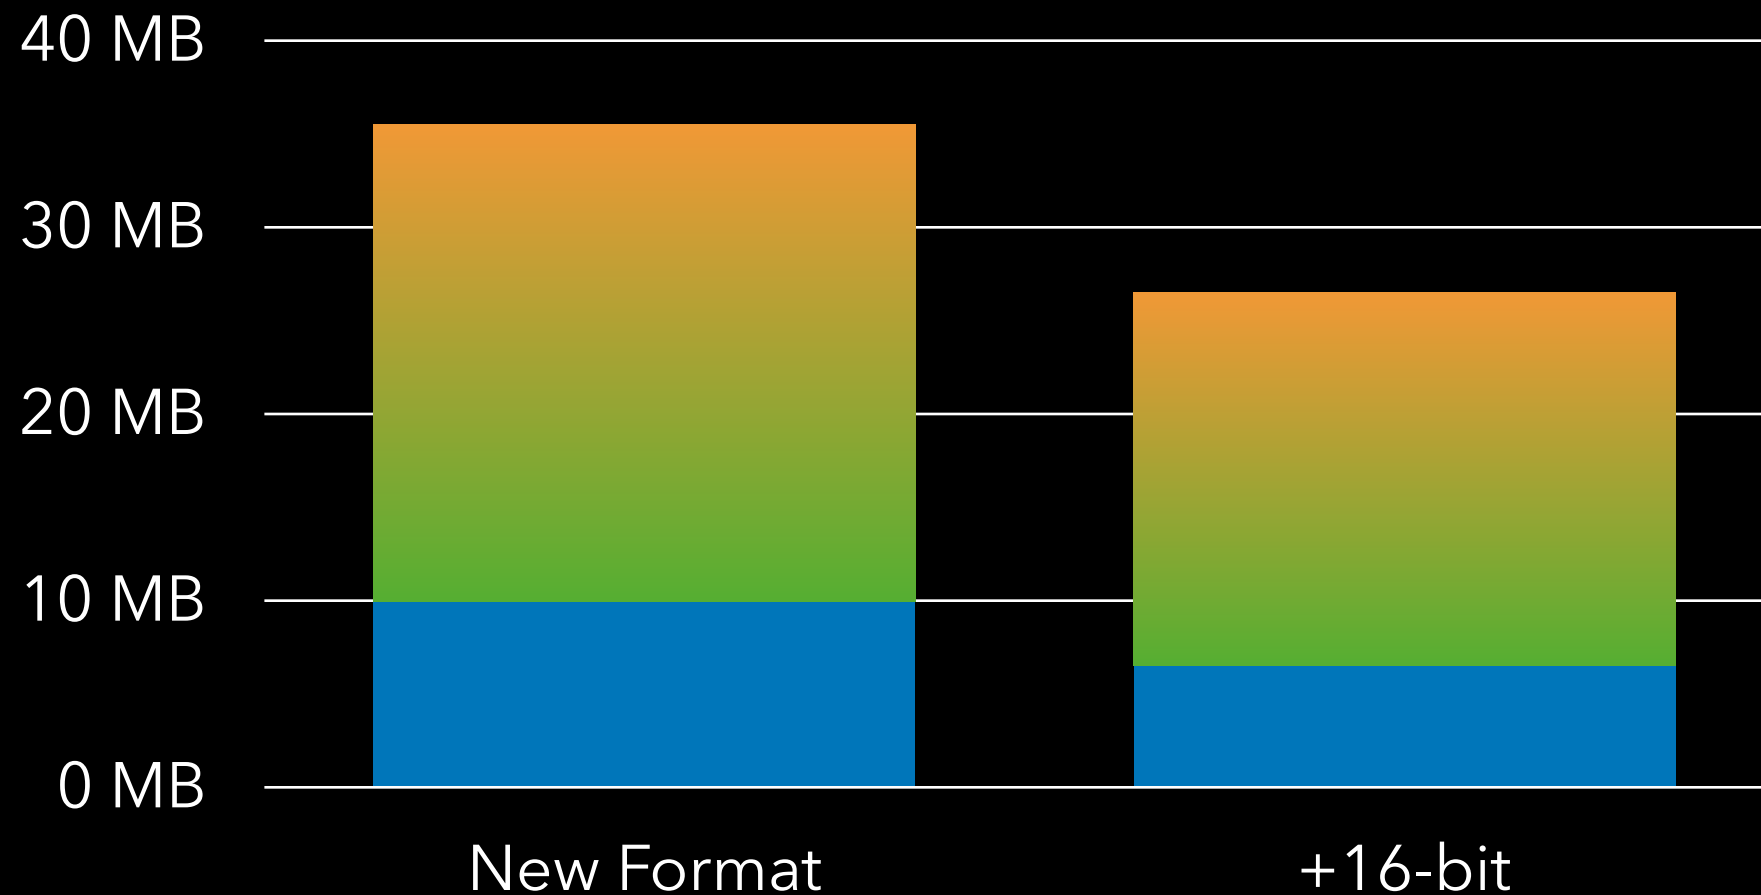
Header			Payload			
<i>0x0</i>	<i>0x2</i>	...	<i>0x80</i>	<i>0x90</i>	<i>0xA0</i>	...
<i>op_add</i> <i>0x80</i>	<i>op_call</i> <i>0xA0</i>	...	OpAdd::Metadata[0]	OpAdd::Metadata[1]	OpCall::Metadata[0]	...

gmail.com



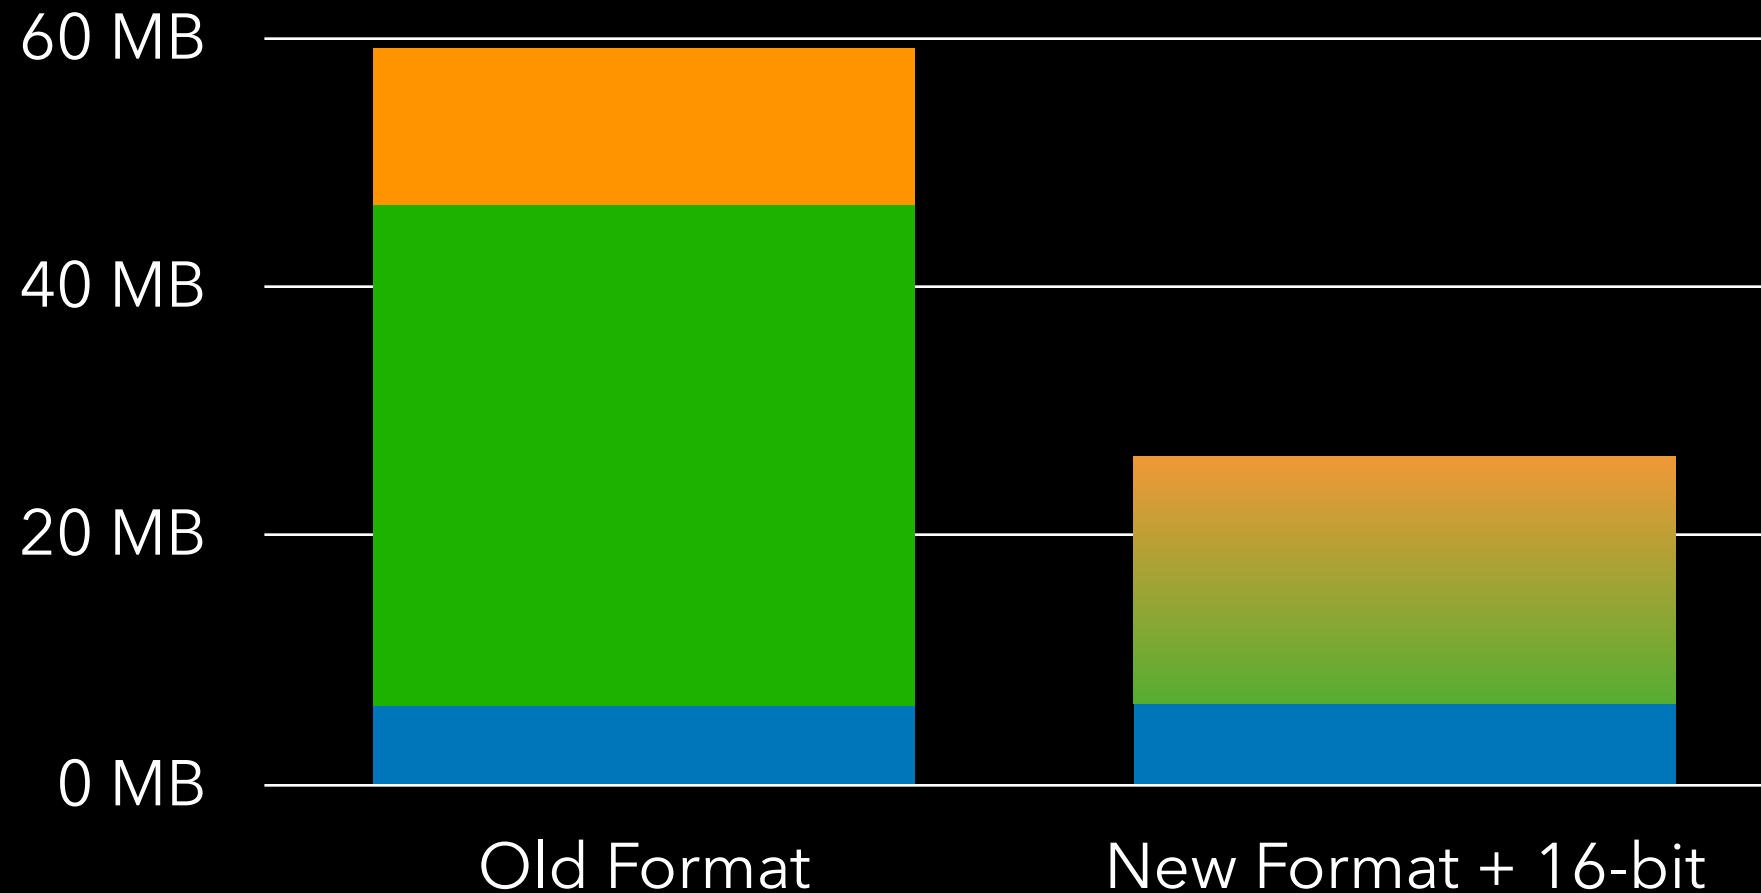
Description		Old Format	New Format	+ 16-bit
	Unlinked	6.17 MB	9.89 MB	6.40 MB
	Linked	40.28 MB	25.51 MB	20.03 MB
	Metadata	12.75 MB		
Total		59.21 MB	35.40 MB	26.42 MB

gmail.com



Description		New Format	+ 16-bit	%
<div></div>	Unlinked	9.89 MB	6.40 MB	-35%
<div></div>	Linked	25.51 MB	20.03 MB	-21%
<div></div>	Metadata			
Total		35.40 MB	26.42 MB	-26%

gmail.com



Description	Before	16-bit	%
Unlinked	6.17 MB	6.40 MB	+4%
Linked	40.28 MB	20.03 MB	-62%
Metadata	12.75 MB		
Total	59.21 MB	26.42 MB	-55%

Agenda

- High level overview
- Old bytecode format
- New bytecode format
- Memory comparison
- Type safety improvements

Old Instruction Definition

```
{ "name": "op_add", "length": 5 }
```


Old Instruction Access

```
SLOW_PATH_DECL(slow_path_add)
{
    JSValue lhs = OP_C(2).jsValue();
    JSValue rhs = OP_C(3).jsValue();
    ...
}
```

Old Instruction Access

```
SLOW_PATH_DECL(slow_path_add)
{
    JSValue lhs = exec->r(pc[2].u.operand).jsValue();
    JSValue rhs = exec->r(pc[3].u.operand).jsValue();
    ...
}
```

Old Instruction Access

```
SLOW_PATH_DECL(slow_path_add)
{
    JSValue lhs = exec->r(pc[2].u.operand).jsValue();
    JSValue rhs = exec->r(pc[3].u.operand).jsValue();
    ...
}
```

Old Instruction Access

```
union {  
    void* pointer;  
    Opcode opcode;  
    int operand;  
    unsigned unsignedValue;  
   WriteBarrierBase<Structure> structure;  
    StructureID structureID;  
   WriteBarrierBase<SymbolTable> symbolTable;  
   WriteBarrierBase<StructureChain> structureChain;  
   WriteBarrierBase<JSCell> jsCell;  
   WriteBarrier<Unknown>* variablePointer;  
    Special::Pointer specialPointer;  
    PropertySlot::GetValueFunc getterFunc;  
    LLIntCallLinkInfo* callLinkInfo;  
    UniquedStringImpl* uid;
```

New Instruction Definition

```
op :add,  
  args: {  
    dst: VirtualRegister,  
    lhs: VirtualRegister,  
    rhs: VirtualRegister,  
    operandTypes: OperandTypes,  
  },  
  metadata: {  
    arithProfile: ArithProfile,  
  }
```

Opcode Struct

```
struct OpAdd : public Instruction {  
    static constexpr OpcodeID opcodeID = op_add;  
  
    VirtualRegister m_dst;  
    VirtualRegister m_lhs;  
    VirtualRegister m_rhs;  
    OperandTypes m_operandTypes;  
    unsigned m_metadataID;  
};
```

Metadata Struct

```
struct OpAdd::Metadata {  
    WTF_MAKE_NONCOPYABLE(Metadata);  
  
public:  
    Metadata(const OpAdd& __op)  
        : m_arithProfile(__op.m_operandTypes)  
    { }  
  
    ArithProfile m_arithProfile;  
};
```

Autogenerate all the things!

- Instruction fitting
- Instruction decoding (narrow vs wide)
- Pretty printing
- Constants for offlineasm
- Opcode IDs
- ...

New Instruction Access

```
SLOW_PATH_DECL(slow_path_add)
{
    OpAdd bytecode = pc->as<OpAdd>();
    JSValue lhs = GET_C(bytecode.m_lhs);
    JSValue rhs = GET_C(bytecode.m_rhs);
    ...
}
```

New Instruction Access

```
SLOW_PATH_DECL(slow_path_add)
{
    OpAdd bytecode = pc->as<OpAdd>();
    JSValue lhs = exec->r(bytecode.m_lhs.offset());
    JSValue rhs = exec->r(bytecode.m_rhs.offset());
    ...
}
```

New Instruction Access

```
SLOW_PATH_DECL(slow_path_add)
{
    OpAdd bytecode = pc->as<OpAdd>();
    JSValue lhs = exec->r(bytecode.m_lhs.offset());
    JSValue rhs = exec->r(bytecode.m_rhs.offset());
    ...
}
```

Thank you!

@tadeuzagallo